

11. Topic Models II: Beyond LDA

DS-GA 1015, Text as Data
Arthur Spirling

April 23, 2019

Housekeeping

- ① HW3 out now. Coming in May 3. Note academic honesty policy.

- ① HW3 out now. Coming in May 3. Note academic honesty policy.
- ② Been called for jury service on May 7, so next week is last lecture (from me at least—Pedro may give embeddings paper/discussion).

- ① HW3 out now. Coming in May 3. Note academic honesty policy.
- ② Been called for jury service on May 7, so next week is last lecture (from me at least—Pedro may give embeddings paper/discussion).

Last time...

Last time...

Topics

gene 0.04
dna 0.02
genetic 0.01
...

life 0.02
evolve 0.01
organism 0.01
...

brain 0.04
neuron 0.02
nerve 0.01
...

data 0.02
number 0.02
computer 0.01
...


Documents

Seeking Life's Bare (Genetic) Necessities

COLD SPRING HARBOR, NEW YORK—How many **genes** does an **organism** need to **survive**? Last week at the genome meeting here,* two genome researchers with radically different approaches presented complementary views of the basic genes needed for **life**. One research team, using **computer** analyses to compare known **genomes**, concluded that today's **organisms** can be sustained with just 250 genes, and that the earliest life forms required a mere 128 **genes**. The other researcher mapped genes in a simple parasite and estimated that for this organism, 800 genes are plenty to do the job—but that anything short of 100 wouldn't be enough.

Although the numbers don't match precisely, those **predictions**

"are not all that far apart," especially in comparison to the 75,000 **genes** in the human genome, notes Siv Andersson of Uppsala University in Sweden, who arrived at the 800 number. But coming up with a consensus answer may be more than just a **genetic** numbers game, particularly as more and more **genomes** are completely mapped and sequenced. "It may be a way of organizing any newly **sequenced genome**," explains Arcady Mushegian, a **computational** molecular biologist at the National Center for Biotechnology Information (NCBI) in Bethesda, Maryland. Comparing an



* Genome Mapping and Sequencing, Cold Spring Harbor, New York, May 8 to 12.

Stripping down. Computer analysis yields an estimate of the minimum modern and ancient genomes.

Topic proportions and assignments



Extensions and Special Cases

Extensions and Special Cases

'vanilla' LDA is a popular model.

Extensions and Special Cases

'vanilla' LDA is a popular model. But assumptions it makes are restrictive/inappropriate for some contexts...

Extensions and Special Cases

'vanilla' LDA is a popular model. But assumptions it makes are restrictive/inappropriate for some contexts...

Assn each document is mix of **multiple** topics.

Extensions and Special Cases

'vanilla' LDA is a popular model. But assumptions it makes are restrictive/inappropriate for some contexts...

Assn each document is mix of **multiple** topics.

But each document is **one** topic. [EAM]

Extensions and Special Cases

'vanilla' LDA is a popular model. But assumptions it makes are restrictive/inappropriate for some contexts...

Assn each document is mix of **multiple** topics.

But each document is **one** topic. [EAM]

Assn topics in documents are **uncorrelated**.

Extensions and Special Cases

'vanilla' LDA is a popular model. But assumptions it makes are restrictive/inappropriate for some contexts...

Assn each document is mix of **multiple** topics.

But each document is **one** topic. [EAM]

Assn topics in documents are **uncorrelated**.

But given topic A , we are more likely to see topic B than C . [CTM]

Extensions and Special Cases

'vanilla' LDA is a popular model. But assumptions it makes are restrictive/inappropriate for some contexts...

Assn each document is mix of **multiple** topics.

But each document is **one** topic. [EAM]

Assn topics in documents are **uncorrelated**.

But given topic A , we are more likely to see topic B than C . [CTM]

Assn documents are **exchangeable** (over time).

Extensions and Special Cases

‘vanilla’ LDA is a popular model. But assumptions it makes are restrictive/inappropriate for some contexts...

Assn each document is mix of **multiple** topics.

But each document is **one** topic. [EAM]

Assn topics in documents are **uncorrelated**.

But given topic A , we are more likely to see topic B than C . [CTM]

Assn documents are **exchangeable** (over time).

But time matters for what topic ‘means’ [DTM]

Extensions and Special Cases

'vanilla' LDA is a popular model. But assumptions it makes are restrictive/inappropriate for some contexts...

Assn each document is mix of **multiple** topics.

But each document is **one** topic. [EAM]

Assn topics in documents are **uncorrelated**.

But given topic A , we are more likely to see topic B than C . [CTM]

Assn documents are **exchangeable** (over time).

But time matters for what topic 'means' [DTM]

Assn no **covariates**

Extensions and Special Cases

'vanilla' LDA is a popular model. But assumptions it makes are restrictive/inappropriate for some contexts...

Assn each document is mix of **multiple** topics.

But each document is **one** topic. [EAM]

Assn topics in documents are **uncorrelated**.

But given topic A , we are more likely to see topic B than C . [CTM]

Assn documents are **exchangeable** (over time).

But time matters for what topic 'means' [DTM]

Assn no **covariates**

But topic prevalence and topic content are $f(X)$ [STM]

Lots of other ideas!

Lots of other ideas!

hierarchical LDA, pachinko allocation, nonparametric pachinko allocation, factorial LDA, gamma-poisson factorization, shared component topic models, dirichlet multinomial regression topic models, nested hierarchical dirichlet process topic model, focused topic model, inverse regression topic model, ideal point topic model, discrete infinite logistic normal topic model multilingual topic model, markov topic model, relational topic model, syntactic topic model, supervised latent dirichlet allocation

Expressed Agenda Model (Grimmer, 2010)

Expressed Agenda Model (Grimmer, 2010)

Suppose each document is assigned to **one** topic.

Expressed Agenda Model (Grimmer, 2010)

Suppose each document is assigned to **one** topic.

Each author allocates a **latent** (to us) proportion of time to each topic.

Expressed Agenda Model (Grimmer, 2010)

Suppose each document is assigned to **one** topic.

Each author allocates a **latent** (to us) proportion of time to each topic.

→ special case of LDA, and used for measuring way Senators 'express' themselves to constituents via **press releases**.

Notice that **set** of topics is **same** across Senators,

Expressed Agenda Model (Grimmer, 2010)

Suppose each document is assigned to **one** topic.


Each author allocates a **latent** (to us) proportion of time to each topic.

→ special case of LDA, and used for measuring way Senators 'express' themselves to constituents via **press releases**.


Notice that **set** of topics is **same** across Senators, but **weights** are allowed to **vary** across Senators.

Senators from same states have similar agendas

Senators from same states have similar agendas



Senators from same states have similar agendas



Senators from same states talk about more similar things than Senators from different states (generally).

Correlated Topic Model (Blei & Lafferty, 2007)

Correlated Topic Model (Blei & Lafferty, 2007)

If a news article talks about `finance` it's more likely to also talk about `law` than it is `baseball`.

Correlated Topic Model (Blei & Lafferty, 2007)

If a news article talks about `finance` it's more likely to also talk about `law` than it is `baseball`. But LDA doesn't allow this (consequence of Dirichlet prior on topic proportions).

Correlated Topic Model (Blei & Lafferty, 2007)

If a news article talks about finance it's more likely to also talk about law than it is baseball. But LDA doesn't allow this (consequence of Dirichlet prior on topic proportions).

The **Correlated Topic Model** allows for **positive covariance** between topics.

Correlated Topic Model (Blei & Lafferty, 2007)

If a news article talks about `finance` it's more likely to also talk about `law` than it is `baseball`. But LDA doesn't allow this (consequence of Dirichlet prior on topic proportions).

The [Correlated Topic Model](#) allows for [positive covariance](#) between topics. Does this by drawing topic proportions from a [log normal](#).

Correlated Topic Model (Blei & Lafferty, 2007)

If a news article talks about `finance` it's more likely to also talk about `law` than it is `baseball`. But LDA doesn't allow this (consequence of Dirichlet prior on topic proportions).

The **Correlated Topic Model** allows for **positive covariance** between topics. Does this by drawing topic proportions from a **log normal**.

Shows improved model fit over LDA.

Correlated Topic Model (Blei & Lafferty, 2007)

If a news article talks about `finance` it's more likely to also talk about `law` than it is `baseball`. But LDA doesn't allow this (consequence of Dirichlet prior on topic proportions).


The **Correlated Topic Model** allows for **positive covariance** between topics. Does this by drawing topic proportions from a **log normal**.

Shows improved model fit over LDA. BTW, note that STM (below) reduces to CTM if no covariates are specified.

Dynamic Topic Model (Blei & Lafferty, 2006)


Dynamic Topic Model (Blei & Lafferty, 2006)

Recall LDA. . .



Dynamic Topic Model (Blei & Lafferty, 2006)


Recall LDA. . .



. . . there are multiple documents, but we don't care about their order.

Dynamic Topic Model (Blei & Lafferty, 2006)


Recall LDA. . .



. . . there are multiple documents, but we don't care about their order.
Our results are the 'same' **regardless** of how we reorder the documents and feed them to the model.

Dynamic Topic Model (Blei & Lafferty, 2006)

Recall LDA. . .




. . . there are multiple documents, but we don't care about their order. Our results are the 'same' **regardless** of how we reorder the documents and feed them to the model.

Dynamic Topic Model has a different model for each time period,

Dynamic Topic Model (Blei & Lafferty, 2006)

Recall LDA...




... there are multiple documents, but we don't care about their order. Our results are the 'same' **regardless** of how we reorder the documents and feed them to the model.


Dynamic Topic Model has a different model for each time period, with topics allowed to evolve over time...

So...

So...



So...



Now, mean parameters for the topic proportions (α s) and what's in the topics (in terms of words, β s) are connected over time via a simple evolutionary process (West & Harrison, 1997).

How to Analyze Political Attention with Minimal Assumptions and Costs (Quinn et al, 2010)

How to Analyze Political Attention with Minimal Assumptions and Costs (Quinn et al, 2010)

What is the agenda of the Senate?

How to Analyze Political Attention with Minimal Assumptions and Costs (Quinn et al, 2010)

What is the **agenda** of the Senate? → dynamic topic model of over 100,000 speeches.

How to Analyze Political Attention with Minimal Assumptions and Costs (Quinn et al, 2010)

What is the **agenda** of the Senate? → dynamic topic model of over 100,000 speeches.

Assume slightly different evolution process (Dynamic Linear Model),

How to Analyze Political Attention with Minimal Assumptions and Costs (Quinn et al, 2010)

What is the **agenda** of the Senate? → dynamic topic model of over 100,000 speeches.

Assume slightly different evolution process (Dynamic Linear Model), and only **one** topic per speech (like Grimmer).

How to Analyze Political Attention with Minimal Assumptions and Costs (Quinn et al, 2010)

What is the **agenda** of the Senate? → dynamic topic model of over 100,000 speeches.

Assume slightly different evolution process (Dynamic Linear Model), and only **one** topic per speech (like Grimmer). Model is fit differently too.


How to Analyze Political Attention with Minimal Assumptions and Costs (Quinn et al, 2010)

What is the **agenda** of the Senate? → dynamic topic model of over 100,000 speeches.

Assume slightly different evolution process (Dynamic Linear Model), and only **one** topic per speech (like Grimmer). Model is fit differently too.

BTW, paper has a lot of **validation**!

Attention to Defense [Use of Force]



(b) *The Number of Words on the 'Defense [Use of Force]' Topic Per Day*

Structural Topic Model (Roberts et al.)

Structural Topic Model (Roberts et al.)

STM = LDA + contextual information.

Structural Topic Model (Roberts et al.)

STM = LDA + contextual information.

→ topic prevalence varies by covariates

Structural Topic Model (Roberts et al.)

STM = LDA + contextual information.

→ topic prevalence varies by covariates

e.g. women may report issues with depression more than men do.

Structural Topic Model (Roberts et al.)

STM = LDA + contextual information.

→ topic prevalence varies by covariates

e.g. women may report issues with depression more than men do.

→ topic content varies by covariates

Structural Topic Model (Roberts et al.)

STM = LDA + contextual information.

→ topic prevalence varies by covariates

e.g. women may report issues with depression more than men do.

→ topic content varies by covariates

e.g. young people ($X = 0$) may talk about depression differently to the way old people ($X = 1$) talk about it.

Structural Topic Model (Roberts et al.)

STM = LDA + contextual information.

→ topic prevalence varies by covariates

e.g. women may report issues with depression more than men do.


→ topic content varies by covariates

e.g. young people ($X = 0$) may talk about depression differently to the way old people ($X = 1$) talk about it.



Including covariates allows for (a) more accurate estimation and (b) better interpretability.

Compare: Plate Diagram

Compare: Plate Diagram



Compare: Plate Diagram



$$\beta \propto \exp(m_v + \kappa_{k,v}^{(r)} + \kappa_{y_d,v}^{(c)} + \kappa_{y_d,k,v}^{(i)})$$


Compare: Per Document Topic Distribution (θ)

Compare: Per Document Topic Distribution (θ)

LDA: each document
has some topic
distribution.

Compare: Per Document Topic Distribution (θ)


LDA: each document has some topic distribution.



Compare: Per Document Topic Distribution (θ)

LDA: each document has some topic distribution.

STM, that topic distribution ('prevalence') is a function of the document metadata.




Compare: Per Document Topic Distribution (θ)

LDA: each document has some topic distribution.


STM, that topic distribution ('prevalence') is a function of the document metadata.

e.g. perhaps male author ($X = 0$) documents have different topics relative to female ($X = 1$) author docs.





Compare: Per Document Topic Distribution (θ)

LDA: each document has some topic distribution.



STM, that topic distribution ('prevalence') is a function of the document metadata.

e.g. perhaps male author ($X = 0$) documents have different topics relative to female ($X = 1$) author docs.




Compare: Per Topic Word Distribution (β)

Compare: Per Topic Word Distribution (β)

LDA: topic ('immigration') has a given distribution over words.


Compare: Per Topic Word Distribution (β)

LDA: topic ('immigration') has a given distribution over words.



Compare: Per Topic Word Distribution (β)

LDA: topic ('immigration') has a given distribution over words.



STM: that word distribution ('content') is a function of the document metadata.

STM: that word distribution ('content') is a function of the document metadata.

e.g. perhaps right parties ($Y = 0$) talk about a given topic differently to left ($Y = 1$) parties.

STM: that word distribution ('content') is a function of the document metadata.

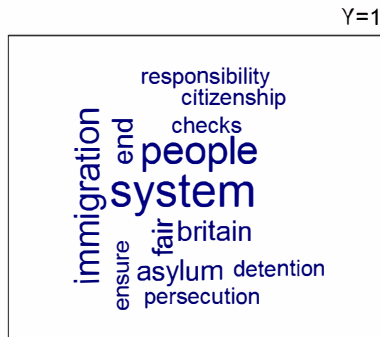
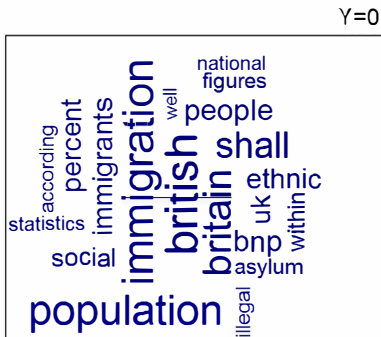
e.g. perhaps right parties ($Y = 0$) talk about a given topic differently to left ($Y = 1$) parties.

In practice, content needs to a single discrete variable.

STM: that word distribution ('content') is a function of the document metadata.

e.g. perhaps right parties ($Y = 0$) talk about a given topic differently to left ($Y = 1$) parties.

In practice, content needs to a single discrete variable.



Partner Exercise

Partner Exercise

- 1 In the dynamic topic model, any changes overtime are 'smooth'. Give an example of a time series problem where you want topics to change over time, but **not** in a smooth way.

Partner Exercise

- 1 In the dynamic topic model, any changes overtime are 'smooth'. Give an example of a time series problem where you want topics to change over time, but **not** in a smooth way.
- 2 In social science, we might be interested in the dynamic topic model because both its α s and the β s are allowed to change. What types of substantive problems do changes in these two different parameters help us model?

Partner Exercise

- 1 In the dynamic topic model, any changes overtime are 'smooth'. Give an example of a time series problem where you want topics to change over time, but **not** in a smooth way.
- 2 In social science, we might be interested in the dynamic topic model because both its α s and the β s are allowed to change. What types of substantive problems do changes in these two different parameters help us model?
- 3 Are the 'effects' in the STM causal? If not, why not, and can you give a scenario where they would be?

Embeddings

Distributional Hypothesis

Distributional Hypothesis

"A word is characterized by the company it keeps."

-Firth, 1957

Distributional Hypothesis

"A word is characterized by the company it keeps."

-Firth, 1957

So: words that are used in the same contexts tend to be similar in their **meaning**.

Distributional Hypothesis

"A word is characterized by the company it keeps."

-Firth, 1957

So: words that are used in the same contexts tend to be similar in their **meaning**. How can we use this insight to understand what words mean?

Distributional Hypothesis

"A word is characterized by the company it keeps."

-Firth, 1957

So: words that are used in the same contexts tend to be similar in their **meaning**. How can we use this insight to understand what words mean?

If the word `liberal` shows up 'near' `privatization` in one country but 'near' `socialist` in another, we can learn something about those polities.

Distributional Hypothesis

"A word is characterized by the company it keeps."

-Firth, 1957

So: words that are used in the same contexts tend to be similar in their **meaning**. How can we use this insight to understand what words mean?

If the word `liberal` shows up 'near' `privatization` in one country but 'near' `socialist` in another, we can learn something about those polities.

We can systematically learn about **analogies** and **similarities**.

Approaches

Approaches

1 'count-based' methods:

Approaches

- 1 'count-based' methods: look at how often words co-occur with neighbors.

Approaches

- 1 'count-based' methods: look at how often words co-occur with neighbors.

e.g. Latent Semantic Analysis: it groups documents using *similar words* and words occurring in *similar documents*.

Approaches

1 'count-based' methods: look at how often words co-occur with neighbors.

e.g. Latent Semantic Analysis: it groups documents using *similar words* and words occurring in *similar documents*. There is some rank reduction to discover latent components.

Approaches

1 'count-based' methods: look at how often words co-occur with neighbors.

e.g. Latent Semantic Analysis: it groups documents using *similar words* and words occurring in *similar documents*. There is some rank reduction to discover latent components.

2 'predictive' methods:

Approaches

1 'count-based' methods: look at how often words co-occur with neighbors.

e.g. Latent Semantic Analysis: it groups documents using *similar words* and words occurring in *similar documents*. There is some rank reduction to discover latent components.

2 'predictive' methods: try to predict a word from knowing its neighbors.

Approaches

- 1 'count-based' methods: look at how often words co-occur with neighbors.

e.g. Latent Semantic Analysis: it groups documents using *similar words* and words occurring in *similar documents*. There is some rank reduction to discover latent components.

- 2 'predictive' methods: try to predict a word from knowing its neighbors. To do this, we *learn* an **embedding vector** that captures meaning of a given word in our corpus a numerical way.

Approaches

- 1 'count-based' methods: look at how often words co-occur with neighbors.

e.g. Latent Semantic Analysis: it groups documents using *similar words* and words occurring in *similar documents*. There is some rank reduction to discover latent components.

- 2 'predictive' methods: try to predict a word from knowing its neighbors. To do this, we *learn* an **embedding vector** that captures meaning of a given word in our corpus a numerical way.

e.g. **Word2Vec**:

Approaches

- 1 'count-based' methods: look at how often words co-occur with neighbors.

e.g. Latent Semantic Analysis: it groups documents using *similar words* and words occurring in *similar documents*. There is some rank reduction to discover latent components.

- 2 'predictive' methods: try to predict a word from knowing its neighbors. To do this, we *learn* an **embedding vector** that captures meaning of a given word in our corpus a numerical way.

e.g. **Word2Vec**: a powerful way to create the word vectors. Comes in two types/models, **Continuous Bag of Words** (CBOW) and **Skip-Gram**.

What's the Difference?

What's the Difference?

- 1 Continuous Bag of Words predicts *target* words from source *context* words.

What's the Difference?

- 1 Continuous Bag of Words predicts *target* words from source *context* words.

So, predicts 'dog' from 'the quick brown fox jumped over []'

What's the Difference?

- 1 **Continuous Bag of Words** predicts *target* words from source *context* words.

So, predicts 'dog' from 'the quick brown fox jumped over []'
→ useful in **small** datasets.

What's the Difference?

- 1 **Continuous Bag of Words** predicts *target* words from source *context* words.

So, predicts 'dog' from 'the quick brown fox jumped over []'
→ useful in **small** datasets.

- 2 **Skip-gram** predicts source *context* words from *target* words

What's the Difference?

- 1 **Continuous Bag of Words** predicts *target* words from source *context* words.

So, predicts 'dog' from 'the quick brown fox jumped over []'
→ useful in **small** datasets.

- 2 **Skip-gram** predicts source *context* words from *target* words

So, predicts 'the quick brown fox jumped over []' from 'dog'.

What's the Difference?

- 1 **Continuous Bag of Words** predicts *target* words from source *context* words.

So, predicts 'dog' from 'the quick brown fox jumped over []'
→ useful in **small** datasets.

- 2 **Skip-gram** predicts source *context* words from *target* words

So, predicts 'the quick brown fox jumped over []' from 'dog'.
→ useful in **large** datasets.

Continuous Bag of Words (CBOW): A Primer

Thinking About 'Context'

Thinking About 'Context'

Thinking About 'Context'

Lots of definitions, but common one is **co-occurrence** within a **window**.

Thinking About 'Context'

Lots of definitions, but common one is **co-occurrence** within a **window**.

Consider “the quick brown fox jumped over the lazy dog”...

Thinking About 'Context'

Lots of definitions, but common one is **co-occurrence** within a **window**.

Consider “the quick brown fox jumped over the lazy dog”...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

Set the **context window** to 2.

Thinking About 'Context'

Lots of definitions, but common one is **co-occurrence** within a **window**.

Consider “the quick brown fox jumped over the lazy dog”...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

Set the **context window** to 2.

Given that context window,

Thinking About 'Context'

Lots of definitions, but common one is **co-occurrence** within a **window**.

Consider “the quick brown fox jumped over the lazy dog”...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

Set the **context window** to 2.

Given that context window, what words co-occur with ‘fox’?

Thinking About 'Context'

Lots of definitions, but common one is **co-occurrence** within a **window**.

Consider “the quick brown fox jumped over the lazy dog”...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

Set the **context window** to 2.

Given that context window, what words co-occur with ‘fox’?

Well...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

Well...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

Well...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

And for 'over'...

Well...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

And for 'over'...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

Well...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

And for 'over'...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

So here,

Well...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

And for 'over'...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

So here, 'quick' co-occurred with 'fox',

Well...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

And for 'over'...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

So here, 'quick' co-occurred with 'fox', but not with 'dog'.

Well...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

And for 'over'...

the	quick	brown	fox	jumped	over	the	lazy	dog
-----	-------	-------	-----	--------	------	-----	------	-----

So here, 'quick' co-occurred with 'fox', but not with 'dog'.

Continuous Bag of Words (CBOW)

Continuous Bag of Words (CBOW)

We are trying to predicting the *target* given the context.

Continuous Bag of Words (CBOW)

We are trying to predicting the *target* given the context.

To keep it simple, consider predicting **one** target word, given **one** context word (“single context word CBOW”)

Continuous Bag of Words (CBOW)

We are trying to predicting the *target* given the context.

To keep it simple, consider predicting **one** target word, given **one** context word (“single context word CBOW”)

So, if our document was “the quick brown fox jumped over the lazy dog”,

Continuous Bag of Words (CBOW)

We are trying to predicting the *target* given the context.

To keep it simple, consider predicting **one** target word, given **one** context word (“single context word CBOW”)

So, if our document was “the quick brown fox jumped over the lazy dog”, we start by using **the** as an input to predict **quick** as an output...

Continuous Bag of Words (CBOW)

We are trying to predicting the *target* given the context.

To keep it simple, consider predicting **one** target word, given **one** context word (“single context word CBOW”)

So, if our document was “the quick brown fox jumped over the lazy dog”, we start by using **the** as an input to predict **quick** as an output...

Then, **the** \leftarrow **quick**

Continuous Bag of Words (CBOW)

We are trying to predicting the *target* given the context.

To keep it simple, consider predicting **one** target word, given **one** context word (“single context word CBOW”)

So, if our document was “the quick brown fox jumped over the lazy dog”, we start by using **the** as an input to predict **quick** as an output...

Then, **the** \leftarrow **quick** and **quick** \rightarrow **brown**

Continuous Bag of Words (CBOW)

We are trying to predicting the *target* given the context.

To keep it simple, consider predicting **one** target word, given **one** context word (“single context word CBOW”)

So, if our document was “the quick brown fox jumped over the lazy dog”, we start by using **the** as an input to predict **quick** as an output...

Then, **the** \leftarrow **quick** and **quick** \rightarrow **brown** We move on to brown.

Continuous Bag of Words (CBOW)

We are trying to predicting the *target* given the context.

To keep it simple, consider predicting **one** target word, given **one** context word (“single context word CBOW”)

So, if our document was “the quick brown fox jumped over the lazy dog”, we start by using **the** as an input to predict **quick** as an output...

Then, **the** \leftarrow **quick** and **quick** \rightarrow **brown** We move on to brown.

Then, **quick** \leftarrow **brown** \rightarrow **fox**, **brown** \leftarrow **fox** \rightarrow **jumped**, etc

Continuous Bag of Words (CBOW)

We are trying to predicting the *target* given the context.

To keep it simple, consider predicting **one** target word, given **one** context word (“single context word CBOW”)

So, if our document was “the quick brown fox jumped over the lazy dog”, we start by using **the** as an input to predict **quick** as an output...

Then, **the** \leftarrow **quick** and **quick** \rightarrow **brown** We move on to brown.

Then, **quick** \leftarrow **brown** \rightarrow **fox**, **brown** \leftarrow **fox** \rightarrow **jumped**, etc

Until **lazy** \rightarrow **dog**,

Continuous Bag of Words (CBOW)

We are trying to predicting the *target* given the context.

To keep it simple, consider predicting **one** target word, given **one** context word (“single context word CBOW”)

So, if our document was “the quick brown fox jumped over the lazy dog”, we start by using **the** as an input to predict **quick** as an output...

Then, **the** \leftarrow **quick** and **quick** \rightarrow **brown** We move on to brown.

Then, **quick** \leftarrow **brown** \rightarrow **fox**, **brown** \leftarrow **fox** \rightarrow **jumped**, etc

Until **lazy** \rightarrow **dog**, and ending with **lazy** \leftarrow **dog** (because we have no other context for this one)

The data

The data

Input	Output	Ob #	the	quick	brown	fox	jumped	over	the	lazy	dog

The data

Input	Output	Ob #	the	quick	brown	fox	jumped	over	the	lazy	dog
the											

The data

Input	Output	Ob #	the	quick	brown	fox	jumped	over	the	lazy	dog
the	quick	1	1	0	0	0	0	0	0	0	0

The data

Input	Output	Ob #	the	quick	brown	fox	jumped	over	the	lazy	dog
the	quick	1	1	0	0	0	0	0	0	0	0
quick	the	2	0	1	0	0	0	0	0	0	0

The data

Input	Output	Ob #	the	quick	brown	fox	jumped	over	the	lazy	dog
the	quick	1	1	0	0	0	0	0	0	0	0
quick	the	2	0	1	0	0	0	0	0	0	0
quick	brown	3	0	1	0	0	0	0	0	0	0

The data

Input	Output	Ob #	the	quick	brown	fox	jumped	over	the	lazy	dog
the	quick	1	1	0	0	0	0	0	0	0	0
quick	the	2	0	1	0	0	0	0	0	0	0
quick	brown	3	0	1	0	0	0	0	0	0	0
brown	quick	4	0	0	1	0	0	0	0	0	0
brown	fox	5	0	0	1	0	0	0	0	0	0
fox	brown	6	0	0	0	1	0	0	0	0	0
fox	jumped	7	0	0	0	1	0	0	0	0	0
jumped	fox	8	0	0	0	0	1	0	0	0	0
jumped	over	9	0	0	0	0	1	0	0	0	0
over	jumped	10	0	0	0	0	0	1	0	0	0
over	the	11	0	0	0	0	0	1	0	0	0
the	over	12	0	0	0	0	0	0	1	0	0
the	lazy	13	0	0	0	0	0	0	1	0	0
lazy	the	14	0	0	0	0	0	0	0	1	0
lazy	dog	15	0	0	0	0	0	0	0	1	0
dog	lazy	16	0	0	0	0	0	0	0	0	1

The data

Input	Output	Ob #	the	quick	brown	fox	jumped	over	the	lazy	dog
the	quick	1	1	0	0	0	0	0	0	0	0
quick	the	2	0	1	0	0	0	0	0	0	0
quick	brown	3	0	1	0	0	0	0	0	0	0
brown	quick	4	0	0	1	0	0	0	0	0	0
brown	fox	5	0	0	1	0	0	0	0	0	0
fox	brown	6	0	0	0	1	0	0	0	0	0
fox	jumped	7	0	0	0	1	0	0	0	0	0
jumped	fox	8	0	0	0	0	1	0	0	0	0
jumped	over	9	0	0	0	0	1	0	0	0	0
over	jumped	10	0	0	0	0	0	1	0	0	0
over	the	11	0	0	0	0	0	1	0	0	0
the	over	12	0	0	0	0	0	0	1	0	0
the	lazy	13	0	0	0	0	0	0	1	0	0
lazy	the	14	0	0	0	0	0	0	0	1	0
lazy	dog	15	0	0	0	0	0	0	0	1	0
dog	lazy	16	0	0	0	0	0	0	0	0	1

'one-hot encoding'

Using the data

Using the data

Each row of this data is a 'one hot encoding' representation of the document.

Using the data

Each row of this data is a 'one hot encoding' representation of the document. Is simply a 'dummy variable' representation.

Using the data

Each row of this data is a 'one hot encoding' representation of the document. Is simply a 'dummy variable' representation.

e.g. observation 6 is $[0, 0, 0, 1, 0, 0, 0, 0, 0]$ since it involves the use of fox at that point in the document.

Using the data

Each row of this data is a 'one hot encoding' representation of the document. Is simply a 'dummy variable' representation.

e.g. observation 6 is $[0, 0, 0, 1, 0, 0, 0, 0, 0]$ since it involves the use of fox at that point in the document.

The binary matrix on the right is the input to a [neural network](#).

Using the data

Each row of this data is a 'one hot encoding' representation of the document. Is simply a 'dummy variable' representation.

e.g. observation 6 is $[0, 0, 0, 1, 0, 0, 0, 0, 0]$ since it involves the use of fox at that point in the document.

The binary matrix on the right is the input to a [neural network](#). Ultimately this will produce a series of [weights](#) which are the [word vector representation](#) of the word.

Using the data

Each row of this data is a 'one hot encoding' representation of the document. Is simply a 'dummy variable' representation.

e.g. observation 6 is $[0, 0, 0, 1, 0, 0, 0, 0, 0]$ since it involves the use of fox at that point in the document.

The binary matrix on the right is the input to a **neural network**. Ultimately this will produce a series of **weights** which are the **word vector representation** of the word.

In practice, neural nets used for the most common models are very simple linear ones, but it doesn't hurt to give a (hand-waving) overview. . .

Neural Nets: A Primer

The Basics

The Basics

We will extract/create **linear combinations of inputs** (the X s) as 'new' features. . .

The Basics

We will extract/create **linear combinations of inputs** (the X s) as 'new' features...

...and model Y as a **non-linear** function of those new features.

The Basics

We will extract/create **linear combinations of inputs** (the X s) as 'new' features...

...and model Y as a **non-linear** function of those new features.

They are known as (artificial) **neural nets** (or neural networks) as they were inspired by the structure of human brains,

The Basics

We will extract/create **linear combinations of inputs** (the X s) as 'new' features...

...and model Y as a **non-linear** function of those new features.

They are known as (artificial) **neural nets** (or neural networks) as they were inspired by the structure of human brains, and attempt to **learn** in way brains do.

The Basics

We will extract/create **linear combinations of inputs** (the X s) as 'new' features...

...and model Y as a **non-linear** function of those new features.

They are known as (artificial) **neural nets** (or neural networks) as they were inspired by the structure of human brains, and attempt to **learn** in way brains do.

They can be made arbitrarily complex,

The Basics

We will extract/create **linear combinations of inputs** (the X s) as 'new' features...

...and model Y as a **non-linear** function of those new features.

They are known as (artificial) **neural nets** (or neural networks) as they were inspired by the structure of human brains, and attempt to **learn** in way brains do.

They can be made arbitrarily complex, such that they can capture almost **any** relationship between Y and the X s... but this comes at a cost:

The Basics

We will extract/create **linear combinations of inputs** (the X s) as 'new' features...

...and model Y as a **non-linear** function of those new features.

They are known as (artificial) **neural nets** (or neural networks) as they were inspired by the structure of human brains, and attempt to **learn** in way brains do.

They can be made arbitrarily complex, such that they can capture almost **any** relationship between Y and the X s... but this comes at a cost:

(a) computational

The Basics

We will extract/create **linear combinations of inputs** (the X s) as 'new' features...

...and model Y as a **non-linear** function of those new features.

They are known as (artificial) **neural nets** (or neural networks) as they were inspired by the structure of human brains, and attempt to **learn** in way brains do.

They can be made arbitrarily complex, such that they can capture almost **any** relationship between Y and the X s... but this comes at a cost:

(a) computational—fitting the model to the data is hard computationally

The Basics

We will extract/create **linear combinations of inputs** (the X s) as 'new' features...

...and model Y as a **non-linear** function of those new features.

They are known as (artificial) **neural nets** (or neural networks) as they were inspired by the structure of human brains, and attempt to **learn** in way brains do.

They can be made arbitrarily complex, such that they can capture almost **any** relationship between Y and the X s... but this comes at a cost:

- (a) computational—fitting the model to the data is hard computationally
- (b) interpretation

The Basics

We will extract/create **linear combinations of inputs** (the X s) as 'new' features...

...and model Y as a **non-linear** function of those new features.

They are known as (artificial) **neural nets** (or neural networks) as they were inspired by the structure of human brains, and attempt to **learn** in way brains do.

They can be made arbitrarily complex, such that they can capture almost **any** relationship between Y and the X s... but this comes at a cost:

- (a) computational—fitting the model to the data is hard computationally
- (b) interpretation—understanding what the models says about the DGP is difficult.

The Set Up

The Set Up

We have p X -variables,

The Set Up

We have p X -variables, which could be counts in a DTM, factors, etc.

The Set Up

We have p X -variables, which could be counts in a DTM, factors, etc.
We want to predict/explain Y .

The Set Up

We have p X -variables, which could be counts in a DTM, factors, etc.
We want to predict/explain Y . As usual, we'll make $\mathbf{y} = f(\mathbf{x}) + \epsilon$.

The Set Up

We have p X -variables, which could be counts in a DTM, factors, etc.
We want to predict/explain Y . As usual, we'll make $\mathbf{y} = f(\mathbf{x}) + \epsilon$.

In a **linear regression** we would set this up as

The Set Up

We have p X -variables, which could be counts in a DTM, factors, etc. We want to predict/explain Y . As usual, we'll make $\mathbf{y} = f(\mathbf{x}) + \epsilon$.

In a **linear regression** we would set this up as

$$\mathbf{y} = \beta' \mathbf{x} + \epsilon$$

The Set Up

We have p X -variables, which could be counts in a DTM, factors, etc. We want to predict/explain Y . As usual, we'll make $\mathbf{y} = f(\mathbf{x}) + \epsilon$.

In a **linear regression** we would set this up as

$$\mathbf{y} = \beta' \mathbf{x} + \epsilon$$

In a **projection pursuit regression** (Freidman & Stuetzle), we make $f(\mathbf{x})$ a sum of **non-linear** (non-parametric) functions of (unidimensional projections) of \mathbf{x} . That is,

$$\mathbf{y} = \sum_{m=1}^M g_m(\beta'_m \mathbf{x}) + \epsilon$$

The Set Up

We have p X -variables, which could be counts in a DTM, factors, etc. We want to predict/explain Y . As usual, we'll make $\mathbf{y} = f(\mathbf{x}) + \epsilon$.

In a **linear regression** we would set this up as

$$\mathbf{y} = \beta' \mathbf{x} + \epsilon$$

In a **projection pursuit regression** (Freidman & Stuetzle), we make $f(\mathbf{x})$ a sum of **non-linear** (non-parametric) functions of (unidimensional projections) of \mathbf{x} . That is,

$$\mathbf{y} = \sum_{m=1}^M g_m(\beta'_m \mathbf{x}) + \epsilon$$

Projection Pursuit Regression

Projection Pursuit Regression

Key component is:

Projection Pursuit Regression

Key component is:

$$\sum_{m=1}^M g_m(\beta'_m \mathbf{x})$$

This is an **additive** model, but not in the original X s.

Projection Pursuit Regression

Key component is:

$$\sum_{m=1}^M g_m(\beta'_m \mathbf{x})$$

This is an **additive** model, but not in the original X s. It's additive in the new/derived features: the β'_j s.

Projection Pursuit Regression

Key component is:

$$\sum_{m=1}^M g_m(\beta'_m \mathbf{x})$$

This is an **additive** model, but not in the original X s. It's additive in the new/derived features: the β'_j s.

Both the g_m and the β_j s are unknown, and will be **estimated** (to optimize model fit).

Projection Pursuit Regression

Key component is:

$$\sum_{m=1}^M g_m(\beta'_m \mathbf{x})$$

This is an **additive** model, but not in the original X s. It's additive in the new/derived features: the β'_j s.

Both the g_m and the β_j s are unknown, and will be **estimated** (to optimize model fit).

To keep things simple, let's suppose that we have two X s and that there will be one function $g \dots$

Example

Example

Suppose the 'true' DGP for Y is

$$Y = \frac{1}{[1 + \exp(-5(\frac{X_1 + X_2}{\sqrt{2}} - 0.5)))]}.$$

Example

Suppose the 'true' DGP for Y is

$$Y = \frac{1}{[1 + \exp(-5(\frac{X_1 + X_2}{\sqrt{2}} - 0.5)))]}.$$


A **linear** function of X s will not do a good job of approximating this!

Example

Suppose the 'true' DGP for Y is

$$Y = \frac{1}{[1 + \exp(-5(\frac{X_1 + X_2}{\sqrt{2}} - 0.5)))]}.$$

A **linear** function of X s will not do a good job of approximating this!



So...

So...

We want need a β'_j that can be 'multiplied' by the X s in a linear way...

So...

We want need a β'_j that can be 'multiplied' by the X s in a linear way...

Here, $\frac{X_1+X_2}{\sqrt{2}}$ is our $\beta'_j x$.

So...

We want need a β'_j that can be 'multiplied' by the X s in a **linear** way...

Here, $\frac{X_1+X_2}{\sqrt{2}}$ is our $\beta'_j x$. Which implies that β_j is $(1/\sqrt{2}, 1/\sqrt{2})$.

So...

We want need a β'_j that can be 'multiplied' by the X s in a linear way...

Here, $\frac{X_1+X_2}{\sqrt{2}}$ is our $\beta'_j x$. Which implies that β_j is $(1/\sqrt{2}, 1/\sqrt{2})$.

Then the dot product $\beta'_j x$

So...

We want need a β'_j that can be 'multiplied' by the X s in a **linear** way...

Here, $\frac{X_1+X_2}{\sqrt{2}}$ is our $\beta'_j x$. Which implies that β_j is $(1/\sqrt{2}, 1/\sqrt{2})$.

Then the dot product $\beta'_j x$ —our **derived** features—enters a non-linear function, g which is

So...

We want need a β'_j that can be 'multiplied' by the X s in a **linear** way...

Here, $\frac{X_1+X_2}{\sqrt{2}}$ is our $\beta'_j x$. Which implies that β_j is $(1/\sqrt{2}, 1/\sqrt{2})$.

Then the dot product $\beta'_j x$ —our **derived** features—enters a non-linear function, g which is

$$\frac{1}{[1 + \exp(-5(\beta'_j x - 0.5))]}$$

Turns out that this idea—non-linear functions of linear combinations of X s—allows essentially infinite **flexibility**:

So...

We want need a β'_j that can be 'multiplied' by the X s in a **linear** way...

Here, $\frac{X_1+X_2}{\sqrt{2}}$ is our $\beta'_j x$. Which implies that β_j is $(1/\sqrt{2}, 1/\sqrt{2})$.

Then the dot product $\beta'_j x$ —our **derived** features—enters a non-linear function, g which is

$$\frac{1}{[1 + \exp(-5(\beta'_j x - 0.5))]}$$

Turns out that this idea—non-linear functions of linear combinations of X s—allows essentially infinite **flexibility**: we can approximate anything we like (some smoothness restrictions in practice).

Example II

Example II

Suppose the 'true' DGP for Y is

$$Y = (\beta_j'x + 0.1) \sin \left(\frac{1}{\frac{\beta_j'x}{3} + 0.1} \right)$$

Example II

Suppose the 'true' DGP for Y is

$$Y = (\beta_j'x + 0.1) \sin \left(\frac{1}{\frac{\beta_j'x}{3} + 0.1} \right)$$

A **linear** function of X s will not do a good job of approximating this!

Example II

Suppose the 'true' DGP for Y is

$$Y = (\beta_j'x + 0.1) \sin \left(\frac{1}{\frac{\beta_j'x}{3} + 0.1} \right)$$

A **linear** function of X s will not do a good job of approximating this!

Note that X_2 has no effect on Y in this case.


Example II

Suppose the 'true' DGP for Y is

$$Y = (\beta_j' x + 0.1) \sin \left(\frac{1}{\frac{\beta_j' x}{3} + 0.1} \right)$$

A **linear** function of X s will not do a good job of approximating this!

Note that X_2 has no effect on Y in this case. So, β_j' is $(1, 0)$.



Neural Nets

Neural Nets

Neural nets are a type of **non-linear** statistical models—which makes them like the **projection pursuit regression**.

Neural Nets

Neural nets are a type of **non-linear** statistical models—which makes them like the **projection pursuit regression**.

Can be used for regression (one Y) or **classification**,

Neural Nets

Neural nets are a type of **non-linear** statistical models—which makes them like the **projection pursuit regression**.

Can be used for regression (one Y) or **classification**, in which case we need to predict Y_1, \dots, Y_K class of a given observation.

Neural Nets

Neural nets are a type of **non-linear** statistical models—which makes them like the **projection pursuit regression**.

Can be used for regression (one Y) or **classification**, in which case we need to predict Y_1, \dots, Y_K class of a given observation.

Infinite variety:

Neural Nets

Neural nets are a type of **non-linear** statistical models—which makes them like the **projection pursuit regression**.

Can be used for regression (one Y) or **classification**, in which case we need to predict Y_1, \dots, Y_K class of a given observation.

Infinite variety: (most?) basic is **single layer perceptron** neural net (AKA “single hidden layer back-propagation network”).

Neural Nets

Neural nets are a type of **non-linear** statistical models—which makes them like the **projection pursuit regression**.

Can be used for regression (one Y) or **classification**, in which case we need to predict Y_1, \dots, Y_K class of a given observation.

Infinite variety: (most?) basic is **single layer perceptron** neural net (AKA “single hidden layer back-propagation network”).

Basic idea remains: linear combinations of the X s, which are pushed through non-linear functions to predict class membership (Y_k).

Neural Nets

Neural nets are a type of **non-linear** statistical models—which makes them like the **projection pursuit regression**.

Can be used for regression (one Y) or **classification**, in which case we need to predict Y_1, \dots, Y_K class of a given observation.

Infinite variety: (most?) basic is **single layer perceptron** neural net (AKA “single hidden layer back-propagation network”).

Basic idea remains: linear combinations of the X s, which are pushed through non-linear functions to predict class membership (Y_k).

Specifically...

Specifically...

Create some **linear** combinations of the X s (βx).

Specifically...

Create some **linear** combinations of the X s (βx).

Create some **non-linear** functions of those linear combinations (something like $\frac{1}{1+\exp(s\beta x)}$) (parametric; cf. PPR). Call those functions Z_m , and are the **derived** features.

Specifically...

Create some **linear** combinations of the X s (βx).

Create some **non-linear** functions of those linear combinations (something like $\frac{1}{1+\exp(s\beta x)}$) (parametric; cf. PPR). Call those functions Z_m , and are the **derived** features.

Create linear functions of the Z to predict Y_k ,

Specifically...

Create some **linear** combinations of the X s (βx).

Create some **non-linear** functions of those linear combinations (something like $\frac{1}{1+\exp(s\beta x)}$) (parametric; cf. PPR). Call those functions Z_m , and are the **derived** features.

Create linear functions of the Z to predict Y_k , of the form

$$T_k = \alpha_0 + \alpha_k Z$$

Specifically...

Create some **linear** combinations of the X s (βx).

Create some **non-linear** functions of those linear combinations (something like $\frac{1}{1+\exp(s\beta x)}$) (parametric; cf. PPR). Call those functions Z_m , and are the **derived** features.

Create linear functions of the Z to predict Y_k , of the form

$$T_k = \alpha_0 + \alpha_k Z$$

Transform those outputs into something useful for a class prediction problem,

Specifically...

Create some **linear** combinations of the X s (βx).

Create some **non-linear** functions of those linear combinations (something like $\frac{1}{1+\exp(s\beta x)}$) (parametric; cf. PPR). Call those functions Z_m , and are the **derived** features.


Create linear functions of the Z to predict Y_k , of the form
 $T_k = \alpha_0 + \alpha_k Z$

Transform those outputs into something useful for a class prediction problem, like the **multinomial logit** transform (called 'softmax' in this literature):

$$g_k(T_k) = \frac{\exp(T_k)}{\sum_{t=1}^K \exp(T_k)}$$


Schematic

Schematic




Back to Embeddings (CBOW)

Schematic of CBOW




Schematic of CBOW




In practice...

In practice...



We have a vocabulary of size V : this is the size of the input 'layer' (the context words) and the size of the output 'layer' (the target words).


In practice...



We have a vocabulary of size V : this is the size of the input 'layer' (the context words) and the size of the output 'layer' (the target words).

→ We want to create **embedding vectors**, which sum up how a given word is related to every other word.

In practice...




We have a vocabulary of size V : this is the size of the input 'layer' (the context words) and the size of the output 'layer' (the target words).

→ We want to create **embedding vectors**, which sum up how a given word is related to every other word.


In practice...

In practice...

Our **inputs** (our one-hot encodings) will be transformed via a **hidden** layer of N 'neurons'. How many neurons?




In practice...



Our **inputs** (our one-hot encodings) will be transformed via a **hidden** layer of N 'neurons'. How many neurons?

More \rightarrow better at representing our inputs (bigger embedding vector).

In practice...




Our **inputs** (our one-hot encodings) will be transformed via a **hidden** layer of N 'neurons'. How many neurons?

More \rightarrow better at representing our inputs (bigger embedding vector).

But that means we need more data.

In practice...



Our **inputs** (our one-hot encodings) will be transformed via a **hidden** layer of N 'neurons'. How many neurons?


More \rightarrow better at representing our inputs (bigger embedding vector).

But that means we need more data.

We go from the input layer to the hidden layer via a weight matrix \mathbf{W}_I of size $V \times N$ —each row represents a context word.


In practice...

In practice...



We go from the hidden layer to the **output** layer via a (different) weight matrix \mathbf{W}_O of size $N \times V$ —each column represents a target word.


In practice...



We go from the hidden layer to the **output** layer via a (different) weight matrix \mathbf{W}_O of size $N \times V$ —each column represents a target word.

Then we'll use softmax to calculate $Pr(\text{word}_k | \text{target})$.


In practice...



We go from the hidden layer to the **output** layer via a (different) weight matrix \mathbf{W}_O of size $N \times V$ —each column represents a target word.

Then we'll use softmax to calculate $Pr(\text{word}_k | \text{target})$. We'll maximize this probability, and that will optimize the weight matrices.

In practice...



We go from the hidden layer to the **output** layer via a (different) weight matrix \mathbf{W}_O of size $N \times V$ —each column represents a target word.

Then we'll use softmax to calculate $Pr(\text{word}_k | \text{target})$. We'll maximize this probability, and that will optimize the weight matrices.

The **embedding vector** for a given word is a **row** of \mathbf{W}_I .

CBOW is a very simple, stripped down NN:

Notes

CBOW is a very simple, stripped down NN: it **doesn't** actually have a non-linear transformation (just linear).

Notes

CBOW is a very simple, stripped down NN: it **doesn't** actually have a non-linear transformation (just linear). But lots of related models and extensions do use non-linear transformations.

Notes

CBOW is a very simple, stripped down NN: it **doesn't** actually have a non-linear transformation (just linear). But lots of related models and extensions do use non-linear transformations.

The **softmax** we want to optimize is the probability of seeing a given (target) word, given the index word we saw:

Notes

CBOW is a very simple, stripped down NN: it **doesn't** actually have a non-linear transformation (just linear). But lots of related models and extensions do use non-linear transformations.

The **softmax** we want to optimize is the probability of seeing a given (target) word, given the index word we saw:

$$\Pr(w_j|w_I) = \frac{\exp(u_j)}{\sum_{j'=1}^V \exp(u_{j'})}$$

CBOW is a very simple, stripped down NN: it **doesn't** actually have a non-linear transformation (just linear). But lots of related models and extensions do use non-linear transformations.

The **softmax** we want to optimize is the probability of seeing a given (target) word, given the index word we saw:

$$\Pr(w_j|w_I) = \frac{\exp(u_j)}{\sum_{j'=1}^V \exp(u_{j'})}$$

u_j is, essentially, the model representation of the word: i.e. the **combination** of the (weighted) input vector for a given word multiplied by the (weighted) output vector.

CBOW is a very simple, stripped down NN: it **doesn't** actually have a non-linear transformation (just linear). But lots of related models and extensions do use non-linear transformations.

The **softmax** we want to optimize is the probability of seeing a given (target) word, given the index word we saw:

$$\Pr(w_j|w_I) = \frac{\exp(u_j)}{\sum_{j'=1}^V \exp(u_{j'})}$$

u_j is, essentially, the model representation of the word: i.e. the **combination** of the (weighted) input vector for a given word multiplied by the (weighted) output vector.

→ probability should be one if the target always follows the context word; zero if it never does, etc.

Embedding Vectors

Embedding Vectors

Once we've optimized,

Embedding Vectors

Once we've optimized, we can extract the word specific vectors from W_I as **embedding** vectors (equal in length to number of neurons).


Embedding Vectors

Once we've optimized, we can extract the word specific vectors from W_I as **embedding** vectors (equal in length to number of neurons). These real valued vectors can be used for **analogies** and related tasks,

Embedding Vectors


Once we've optimized, we can extract the word specific vectors from W_I as **embedding** vectors (equal in length to number of neurons). These real valued vectors can be used for **analogies** and related tasks,

$$v_{\text{queen}} - v_{\text{woman}} + v_{\text{men}} \approx v_{\text{king}}$$




Related Tasks

Related Tasks



Verb tense



Country-Capital

Partner Exercise

Partner Exercise

- 1 How do we know whether a word embedding vector is a good representation or not? How could we test the merits of one particular model versus another?
- 2 Embeddings reflect cultural biases. How would you show this in practice given what we discussed above?

Skip-Gram

CBOW can be generalized to **multiple context words**.

Skip-Gram

CBOW can be generalized to **multiple context words**. But **Skip-Gram** generally more popular.

Skip-Gram

CBOW can be generalized to **multiple context words**. But **Skip-Gram** generally more popular.

The **Skip-gram** model predicts the context given the word. This means the output (the context) is **two** hot-encoded vectors (assuming two word context).

Skip-Gram

CBOW can be generalized to **multiple context words**. But **Skip-Gram** generally more popular.

The **Skip-gram** model predicts the context given the word. This means the output (the context) is **two** hot-encoded vectors (assuming two word context).

→ Skip-gram is more subtle:

Skip-Gram

CBOW can be generalized to **multiple context words**. But **Skip-Gram** generally more popular.

The **Skip-gram** model predicts the context given the word. This means the output (the context) is **two** hot-encoded vectors (assuming two word context).

→ Skip-gram is more subtle: contexts can be different for **homographs** (like 'JFK' the airport vs 'JFK' the President).

Skip-Gram

CBOW can be generalized to **multiple context words**. But **Skip-Gram** generally more popular.

The **Skip-gram** model predicts the context given the word. This means the output (the context) is **two** hot-encoded vectors (assuming two word context).

→ Skip-gram is more subtle: contexts can be different for **homographs** (like 'JFK' the airport vs 'JFK' the President).

Shows excellent performance on many tasks (better than CBOW).

Notes II

Notes II

Why does Word2Vec work?

Notes II

Why does Word2Vec work? Area of some interest: see e.g. Levy & Goldberg in NIPS, 2014.

Notes II

Why does Word2Vec work? Area of some interest: see e.g. Levy & Goldberg in NIPS, 2014.

Producing word embedding vectors is computationally expensive, and you need a lot of data.

Notes II

Why does Word2Vec work? Area of some interest: see e.g. Levy & Goldberg in NIPS, 2014.

Producing word embedding vectors is computationally expensive, and you need a lot of data. So **pre-trained** embeddings (esp from Wikipedia) are available.

Notes II

Why does Word2Vec work? Area of some interest: see e.g. Levy & Goldberg in NIPS, 2014.

Producing word embedding vectors is computationally expensive, and you need a lot of data. So **pre-trained** embeddings (esp from Wikipedia) are available.

GloVe, Global **V**ectors for Word Representation, is method that works via dimension reduction of **counts** of co-occurrences (rather than trying to predict, as Word2Vec)

Notes II

Why does Word2Vec work? Area of some interest: see e.g. Levy & Goldberg in NIPS, 2014.

Producing word embedding vectors is computationally expensive, and you need a lot of data. So **pre-trained** embeddings (esp from Wikipedia) are available.

GloVe, Global **V**ectors for Word Representation, is method that works via dimension reduction of **counts** of co-occurrences (rather than trying to predict, as Word2Vec)

Can we provide framework for embeddings so we can talk about one word being statistically significantly different to another? (yes! Cho et al, 2018). And perhaps make embeddings dependent on **covariates**? (Rudolph et al, 2017)